Is HDF5 a good format to replace UVFITS?

Danny C. Price[1], Benjamin R. Barsdell[1], and Lincoln J. Greenhill[1]

[1]*Harvard-Smithsonian Center for Astrophysics, MS 42 60 Garden Street, Cambridge MA 02143*

Abstract. The FITS (Flexible Image Transport System) data format was developed in the late 1970s for storage and exchange of astronomy-related image data. Since then, it has become a standard file format not only for images, but also for radio interferometer data (e.g. UVFITS, FITS-IDI). But is FITS the right format for next-generation telescopes to adopt? The newer Hierarchical Data Format (HDF5) file format offers considerable advantages over FITS, but has yet to gain widespread adoption within radio astronomy. One of the major holdbacks is that HDF5 is not well supported by data reduction software packages. Here, we present a comparison of FITS, HDF5, and the MeasurementSet (MS) format for storage of interferometric data. In addition, we present a tool for converting between formats. We show that the underlying data model of FITS can be ported to HDF5, a first step toward achieving wider HDF5 support.

1.   Introduction

The Flexible Image Transport System (FITS) data format is the most widespread standard for the storage and exchange of datasets within astronomy. Since its inception (Wells & Greisen 1979; Greisen et al. 1980), FITS has enjoyed several decades of widespread usage. The success of FITS has been attributed in part to the guiding maxim "once FITS, always FITS": that changes to the standard must be incremental and must not break backward compability. For this reason, it is familiar to many generations of astronomers, and a large ecosystem of software has in turn motivated further adoption of the standard. In particular, the `CFITSIO` library (ascl:1010.001) for the reading and writing of FITS files has become the *de facto* standard.

FITS has necessarily evolved over the years, with the addition of features such as random groups (Greisen & Harten 1981), tables (Harten et al. 1988; Cotton et al. 1995), and compression (Pence 2002); FITS is now officially at version 3.0 (Pence et al. 2010). However, these changes have been relatively minor iterations upon the core FITS format.

The limitiations of FITS have been previously acknowledged and documented, for example in Thomas et al. (2014a) and Thomas et al. (2014b). In Kitaeff et al. (2014), the authors consider JPEG2000 as an alternative format for images and data cubes. Thomas et al. (2001) discuss advantages of converting FITS files to XML; Jennings et al. (1995) considered HDF4 as a format. Here, we consider the immediate, practicable advantages of HDF5 (Hierarchical Data Format) as a data storage format for visibility data. We show that data in FITS format can be converted in a straightforward fashion to HDF5 format, and that conventions for the storage of visibility data can be ported to HDF5.

1.1. Definitions

In order to discuss data storage methods and file formats without ambiguity, we first need to clarify our vocabulary:

- *Data model:* a high-level, conceptual model of data, types of data, and how data are organized, e.g. "group" and "dataset".

- *Data schema:* a lower-level, domain-specific ontology (i.e. framework that gives meaning) of how data and metadata are arranged inside a data model.

- *Storage model:* how objects from the data model are mapped to bytes within an address space on storage media.

- *Convention:* a documented data schema that has widespread acceptance within a community of users.

- *Standard:* the acknowledged, formal specification of a file format. A standard may or may not define acceptable data models and schema.

From this view, the data model can be seen as *syntax*, while the data schema may be seen as the ontology that gives *semantics*. Neither the FITS nor HDF5 standards define data schema; however there are registered FITS conventions[1] for certain classes of data.

2. Conventions for visibility data storage

2.1. FITS-IDI and UVFITS

There are two registered FITS conventions for the storage of visibility data from synthesis imaging radio telescopes: FITS-IDI (Griesen 2008), and UVFITS (Greisen 2012). Both these formats store not only the visiblility data, but also metadata such as antenna positions, information on observation setup, and calibration tables.

The two visibility data conventions share many keywords and unit definitions, but their schema and underlying storage models differ. In UVFITS, the visibility data are stored in a random group HDU (header data unit), whereas in FITS-IDI data are stored in a binary table HDU. In both formats, each row of the table contains columns for the timestamp and a baseline identifier, along with the multidimensional visibility array for that timestamp and baseline.

2.2. CASA MeasurementSets

An alternative file format for visibility data is the MeasurementSet (MS Kemball & Wieringa 2001). This format is used by the CASA reduction package (ascl:1107.013) and related software. The storage model for the MS format is a directory consisting of several data files nested inside child directories. Unlike FITS and HDF5, MS has no in-built data compression capability. Visibility data is stored in the MAIN table; this also provides keys to the various subtables that provide instrument and observation metadata. The MS standard defines data schema for images, visibility data and single-dish data. The MS data model differ significantly to FITS, meaning that conversion

---

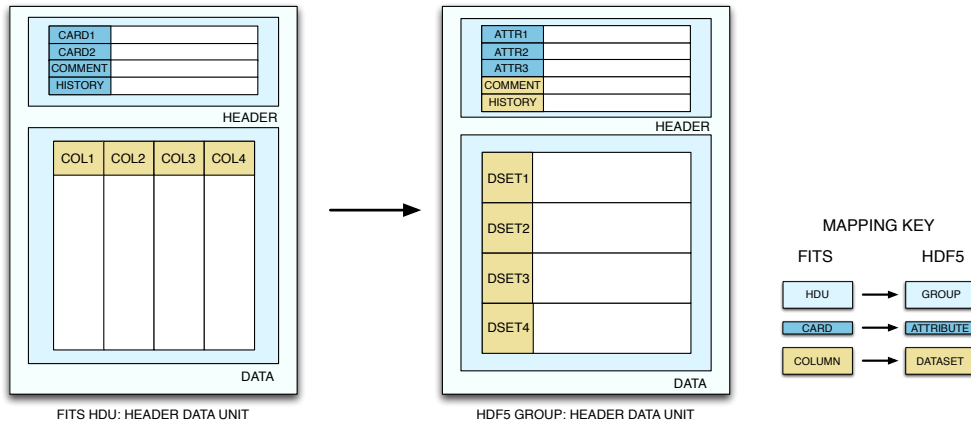[1]`http://fits.gsfc.nasa.gov/fits_registry.html`

Figure 1. Diagram showing the mapping of the FITS data structures into the HDF5 data model.

requires a non-trivial mapping. For example, MS files have no equivalent to the FITS HDU, and while FITS-IDI assigns an integer ID for every baseline, MS uses a pair of antenna IDs.

2.3. HDF5 for visibility storage

There are several advantages of the HDF5 format over both FITS and MS, the most compelling of these are improvements to the storage model that are beneficial for large datasets. For example, HDF5 provides parallel and network I/O, data chunking methods, external (i.e. distributed) object storage, and a filter pipeline for data compression. Of specific interest for visibility data is `bitshuffle`[2], an HDF5 filter designed for fast compression of visibility data. Using `bitshuffle` on a 1.2 GB test dataset of data from the LEDA correlator (Kocz et al. 2014), we achieved lossless compression ratio of 1.65x, with total file compression and write time of 7.5 s; in comparison the data compressed by 1.40x in 53.0 s using standard `gzip`.

HDF5 is already in limited use within astronomy, but no general convention for visibility data storage currently exists. The LOFAR radio telescope has developed conventions specifically for storage of LOFAR data (Wise et al. 2011; Alexov et al. 2012), and the (unfunded) AstroHDF effort sought to explore HDF5 for astronomical datasets further (Masters et al. 2012).

3. Conversion between FITS and HDF5

While there are many possible mappings, one of the simplest is to map the FITS HDU structure to a HDF5 *group*, FITS cards to HDF5 *attributes*, and the FITS data payloads to HDF5 *datasets* (Fig 1). A similar approach is being undetaken to port the N-dimensional data format (NDF) to HDF5 (Jenness et al. 2014). The hierarchical nature of HDF5 allows the header unit and data unit within the FITS HDU to be mapped

---

[2] https://github.com/kiyo-masui/bitshuffle

to HDF5 groups within a parent group; similarly, in the case of FITS tables, each column can be mapped to a separate dataset within a parent group. All FITS data types (e.g. float32) have HDF5 equivalents.

We have implemented a utility called `fits2hdf` that uses this mapping to convert FITS files into HDF5 format[3]. This utility is written in Python, and uses the PyFITS (ascl:1207.009) and h5py libraries for file I/O. These "HDFITS" files have an additional attribute in the root to identify them as having a HDFITS data model. While `fits2hdf` was designed to port UVFITS/FITS-IDI data into HDF5, any valid FITS file may be converted by this utility. As the HDFITS data model is a restricted subset of the complete HDF5 data model, any HDFITS file may be converted back into a FITS file without complication.

A similar approach can be taken with MS files, converting them into "HDF-MS" data models within HDF5. This functionality is provided by a separate utility, `ms2hdf`, provided as part of the `fits2hdf` package.

4. Conclusions

By porting the FITS and MS data models to HDF5, we may begin to leverage the advantages of HDF5. Doing so maintains familiarity to users and allows conversion back into FITS/MS. This also limits the amount of modification required for existing programs to read, write and understand HDF5-stored data.

References

Alexov, A., et al. 2012, in ADASS XXI, 283
Cotton, W. D., Tody, D., & Pence, W. D. 1995, A&A Supplement, 113, 159
Greisen, E. W. 2012, AIPS Memorandum Series, 117, 1
Greisen, E. W., & Harten, R. H. 1981, A&A Supplement, 44, 371
Greisen, E. W., Wells, D. C., & Harten, R. H. 1980, Applications of Digital Image Processing to Astronomy, 264, 298
Griesen, E. W. 2008, AIPS Memorandum Series, 114
Harten, R. H., et al. 1988, A&A Supplement, 73, 365
Jenness, T., et al. 2014, Astronomy and Computing, 1
Jennings, D. G., Pence, W. D., & Folk, M. 1995, ADASS IV, 77, 229
Kemball, A. J., & Wieringa, M. H. 2001, CASA Memo Series
Kitaeff, V. V., Cannon, A., Wicenec, A., & Taubman, D. 2014, arXiv.org. `1403.2801v4`
Kocz, J., et al. 2014, Journal of Astronomical Instrumentation, 03, 1450002
Masters, J., et al. 2012, in ADASS XXI, 871
Pence, W. D. 2002, Astronomical Data Analysis II. Edited by Starck, 4847, 444
Pence, W. D., et al. 2010, A&A, 524, 42
Thomas, B., Shaya, E., & Cheung, C. 2001, ADASS X, 238, 487
Thomas, B., et al. 2014a, in ADASS XXIII, 351
— 2014b, Astronomy and Computing, 1
Wells, D. C., & Greisen, E. W. 1979, in 5th Workshop on Image Processing in Astronomy, 445
Wise, M., et al. 2011, in ADASS XX, 663

---

[3]`https://github.com/telegraphic/fits2hdf`